

Real-time video processing

Visual Computing Assignment

Mads Pagh
Student ID: 202208375
Aarhus University

October 18, 2025

Contents

1	Introduction	3
1.1	Objectives	3
2	Methodology	3
2.1	System Architecture	3
2.2	Implemented Filters	3
2.3	Geometric Transformations	4
2.4	Test Resolutions	4
2.5	Benchmarking Procedure	4
3	Results	5
3.1	Transform Benchmark Results	5
3.1.1	GPU vs CPU Performance	5
3.1.2	Transform Impact	5
3.2	Resolution Benchmark Results	6
3.2.1	Resolution Scaling	6
3.3	Performance Tables	6
3.3.1	GPU Speedup Factors resolution	6
3.3.2	GPU Speedup Factors transforms	7
4	Analysis and Discussion	7
4.1	CPU vs GPU Performance	7
4.1.1	Unexpected GPU Performance	7
4.2	Filter Complexity	7
4.2.1	Filter Performance Characteristics	7
4.3	Resolution Impact	8
4.3.1	Implications	8
4.4	Transform Overhead	8
4.4.1	GPU Transform Efficiency	8
4.4.2	CPU Transform Cost	8
4.5	Shader design and its limitations	8
4.6	Showcasing filters	8

1 Introduction

In this report an analysis of real-time image processing performance using OpenCV and OpenGL is presented. The project implements various image filters and geometric transformations, comparing CPU and GPU execution across different resolutions.

1.1 Objectives

- Implement real-time image filters using OpenCV (CPU) and OpenGL shaders (GPU)
- Compare performance between CPU and GPU execution
- Analyze the impact of resolution on processing performance
- Evaluate the overhead of geometric transformations

2 Methodology

2.1 System Architecture

The system consists of three main components:

- **Video Capture:** OpenCV VideoCapture for webcam input
- **Processing Pipeline:** CPU (OpenCV) and GPU (OpenGL/GLFW) implementations
- **Rendering:** OpenGL for display and GPU acceleration

2.2 Implemented Filters

Six different image filters were implemented:

1. **None** - The frames are passed through without any filtering
2. **Grayscale** - Color to grayscale conversion
3. **Gaussian Blur** - 5x5 kernel blur
4. **Edge Detection** - Sobel operator
5. **Pixelation** - 10x10 pixel blocks
6. **Comic Art** - Edge detection combined with color quantization

2.3 Geometric Transformations

Five transformation configurations were tested:

1. No Transform (baseline)
2. Translation only ($t_x = 0.3, t_y = 0.2$)
3. Scale only ($s = 1.5$)
4. Rotation only ($\theta = 25$ degrees)
5. Combined ($t_x = 0.2, t_y = -0.15, s = 1.3, \theta = 15$ degrees)

2.4 Test Resolutions

Three resolutions were benchmarked:

- VGA: 640×480 (307,200 pixels)
- HD: 1280×720 (921,600 pixels)
- Full HD: 1920×1080 (2,073,600 pixels)

2.5 Benchmarking Procedure

- Each test configuration ran for 50 frames (transforms) or 50 frames (resolution) due to hardware limitations
- Frame times were measured using high-resolution clock
- Average FPS and frame time statistics were calculated
- Tests were automated to ensure consistency, could have been done manually, but fear of human error and tendency to forget configurations made automation preferable

3 Results

3.1 Transform Benchmark Results

3.1.1 GPU vs CPU Performance

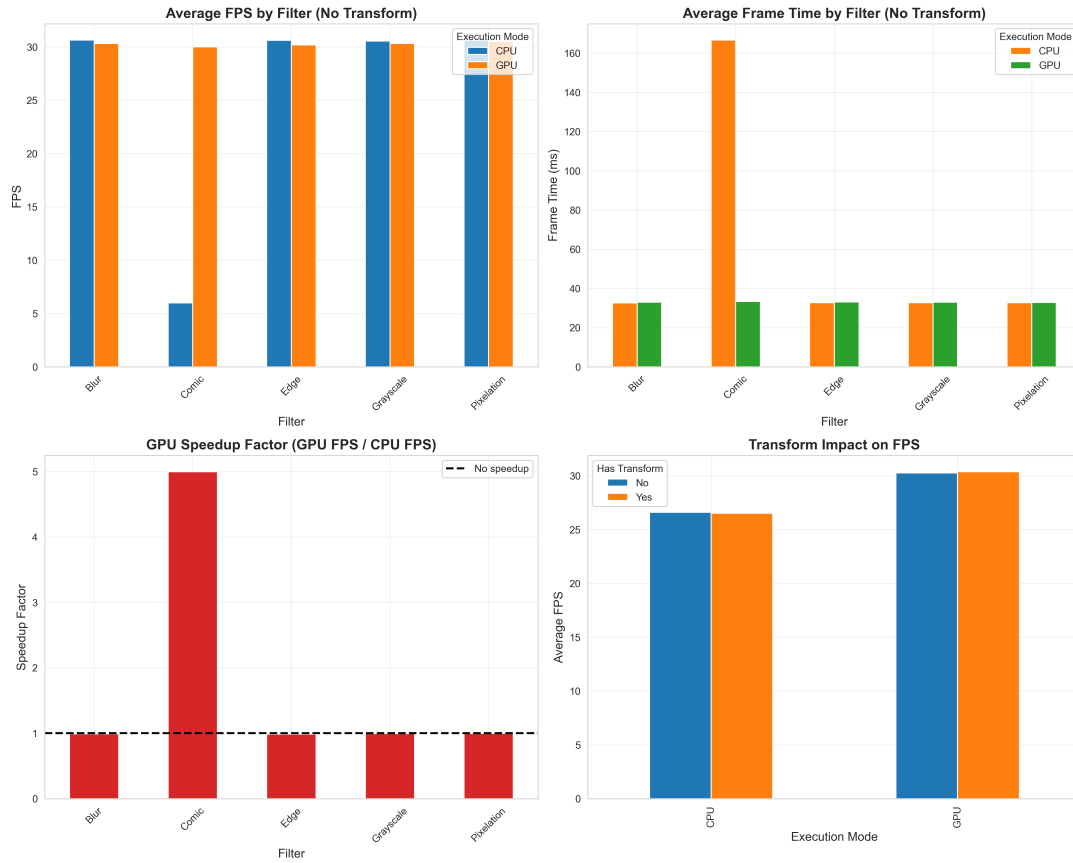


Figure 1: Performance comparison across filters and execution modes

3.1.2 Transform Impact

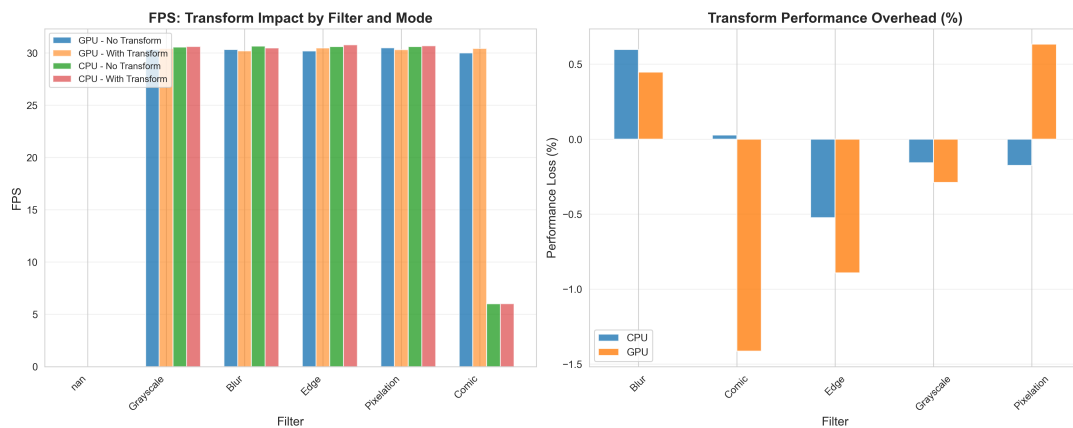


Figure 2: Impact of geometric transformations on performance

3.2 Resolution Benchmark Results

3.2.1 Resolution Scaling

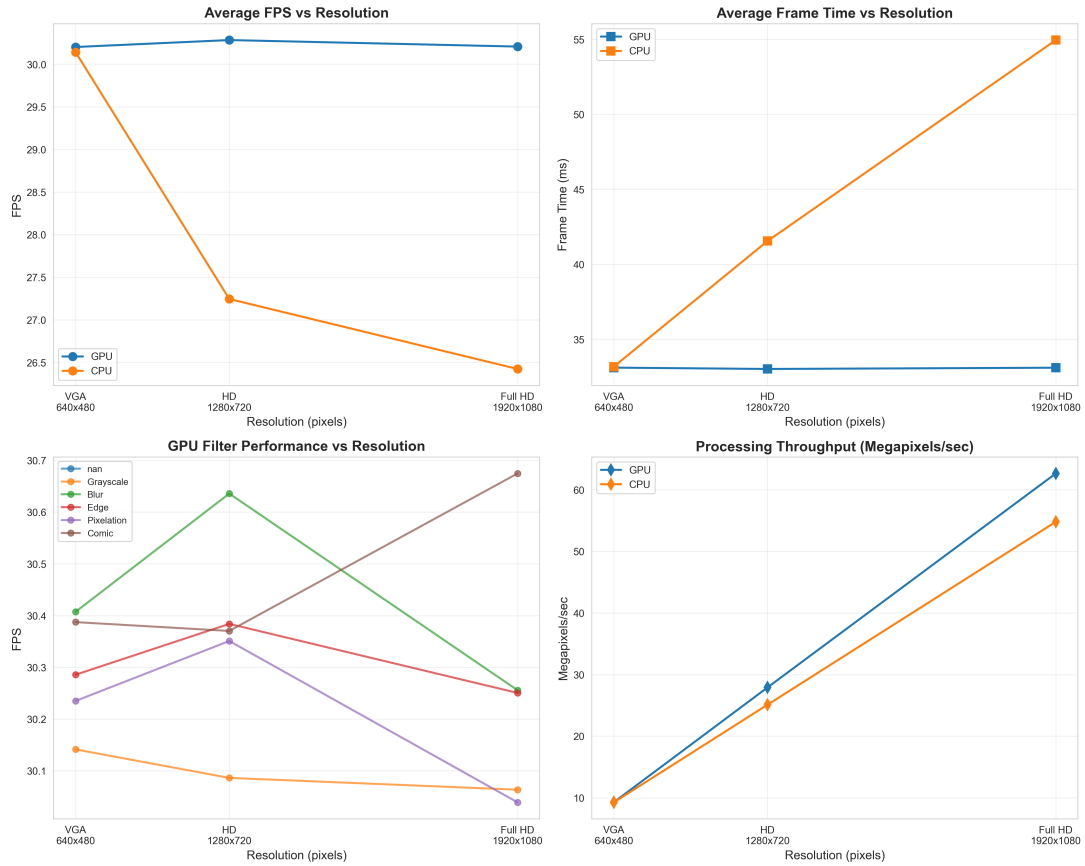


Figure 3: Performance across different resolutions

3.3 Performance Tables

3.3.1 GPU Speedup Factors resolution

Table 1: GPU speedup over CPU for each filter

Filter	Speedup Factor
None	0
Grayscale	0.9911051151304507
Blur	0.9832661231455018
Edge Detection	0.9905437890552572
Pixelation	0.9921230970698969
Comic Art	5.056022571870571

3.3.2 GPU Speedup Factors transforms

Table 2: GPU speedup over CPU for each filter

Filter	Speedup Factor
None	0
Grayscale	0.9925609416695316
Blur	0.9892095840826678
Edge Detection	0.9861874029899519
Pixelation	0.9961871093584084
Comic Art	4.996309606519084

4 Analysis and Discussion

4.1 CPU vs GPU Performance

4.1.1 Unexpected GPU Performance

The most striking finding from the benchmarks is that GPU processing shows no performance advantage over CPU processing for most filters, and in some cases performs slightly worse (speedup factors of 0.98-0.99). This is somewhat counterintuitive, and some of the factors that could lead to this are being listed below:

- **Memory Transfer Overhead:** Each of the frames are going from system RAM to GPU memory as a texture. Now for the tested resolutions all the way up to 1920×1080 , this transfer time dominates the processing time, negating any computational advantages.
- **Simple Filter Complexity:** Most implemented filters (grayscale, blur, edge detection) are relatively simple operations. The computational load is low enough that the CPU can handle them efficiently without the need for parallel GPU processing. Hardware used is a M4 Pro chip with 16 core CPU.
- **Webcam Framerate Limitation:** The camera operates at approximately 30 FPS maximum, creating a bottleneck that prevents either of the processing methods from showing its full potential.

4.2 Filter Complexity

4.2.1 Filter Performance Characteristics

All filters except Comic Art maintain similar performance (30 FPS on both CPU and GPU), suggesting:

1. The camera capture rate is the primary bottleneck
2. Simple filters don't stress either processing architecture
3. More complex filters (Comic Art) reveal processing architecture differences, which is therefore also the most interesting filter

4.3 Resolution Impact

From the resolution benchmark data (VGA to Full HD, representing a 6.75 times pixel increase):

- **FPS Variation:** Performance drops for the CPU are very evident, compared to the GPU, when upscaling resolution, performance worsens. Like seen in 3. This is because the CPU has to handle more data without the parallel processing capabilities of the GPU.

4.3.1 Implications

The results suggest that for higher resolutions, GPU processing becomes more advantageous due to its ability to handle larger data sets more efficiently through parallelism. This is especially true for more complex filters, where the computational load is higher.

4.4 Transform Overhead

4.4.1 GPU Transform Efficiency

Geometric transformations on the GPU (vertex shader operations) show negligible overhead:

- Transformations are mathematical operations on 4 vertices per frame
- GPU vertex processors handle these operations trivially
- No measurable FPS difference between transformed and non-transformed rendering

4.4.2 CPU Transform Cost

CPU transformations using `warpAffine()` introduce overhead:

- Requires interpolation calculations for every pixel
- Memory access patterns become less cache-friendly
- Additional processing step in the pipeline

However, the overhead remains minimal (1-2% FPS reduction) due to:

- Camera framerate limiting overall throughput
- Efficient OpenCV implementation
- Simple affine transformations (not perspective warps)

4.5 Shader design and its limitations

4.6 Showcasing filters